

Runtime verification

(aka Dynamic Analysis)

Yuri Gribov, Samsung Advanced Institute of Technology

[Lalambda '21](#)

About the author

- Teamlead @ Samsung Moscow
- Accidentally became a compiler writer 15 years ago
 - In-house, GCC, LLVM, neurocompilers (also some HPC and gamedev)
- Passionate about verification in general and dynamic/static analyses in particular
 - GitHub [.yugr](#)
 - Habr [the_real_yugr](#)



Disclaimer

- A big picture overview without delving into details of particular checkers or technologies
- Engineering focus
- C/C++-focused (although ideas are generally applicable)

This presentation is available at

- (slides) <https://github.com/yugr/Lambda/blob/master/21/talk.md>
- (practice)
<https://github.com/yugr/Lambda/blob/master/21/practice.md>

Overview

- Runtime verification aka dynamic analysis
- Instrumentation of programs to verify behavioral invariants at runtime
 - safety, performance, etc.
 - verifying code is called a "monitor"
- (Much) more widely used in industry than static tools:
 - no false positives
 - no scalability problems
 - reprocases easily available

Disadvantages

- Limited coverage
- Solved via
 - fuzzing
 - rule/grammar-based input generators
 - A/B testing (in production environments like Google services)

Example analyses

- Virtual memory :)
- Sanity checks in code
 - e.g. C/C++ assertions in programs
 - e.g. Glibc `malloc` or libstdc++ iterator internal checks
- Valgrind
- Sanitizers (Asan, UBSan, Msan, Tsan, etc.)
- "Business rules" (GDPR, data minimization, etc.)

Community

- Academia ([Runtime Verification conference](#))
 - grew out of model checking in 2000-s ([Runtime Verification - 17 Years Later](#))
 - verify complex modal logic formulas on program traces
 - usually applied to interesting but niche projects
- Industry (hackers and corporations)
 - automatically detect bugs at large scale (without manual work by user)
 - much older (malloc debuggers existed at least since 80-s)
 - typically much more influential

Dynamic analysis algorithm

```
errors = {}  
program_with_monitor = instrument(program, spec)  
while test_corpus not empty:  
    test_input = test_corpus.pop()  
    errors, coverage, ... += program_with_monitor(test_input)  
    update test_corpus
```


Dynamic analysis algorithm

```
errors = {}  
program_with_monitor = >>>instrument<<<(program, >>>spec<<<)  
while test_corpus not empty:  
    test_input = test_corpus.pop()  
    errors, coverage, ... += program_with_monitor(test_input)  
>>>update test_corpus<<<
```

Ontology of dynamic analysis project

Runtime analysis tool ("checker") contains of three main "parts":

- spec: an invariant that we want to check
- instrumentation (aka monitor): a way to verify that invariant is preserved during execution
- test corpus: input data which we run the checker through

New successful checkers are created by innovating in any of the three components.

Creating new checkers: spec

The "spec" part:

- come up with a new interesting class of bugs and propose method to autodetect them
- most interesting classes already handled :(
- e.g. [Sortchecker](#) was the first tool to check [qsort axioms](#)

Spec taxonomy (1)

- Memory errors ([Asan](#)/[Msan](#), [Valgrind](#)):
 - liveness errors: accessing after end-of-life (use-after-free, use-after-return, iterator invalidation)
 - buffer overflow: heap, global, stack
 - uninitialized memory
 - memory leaks
- Typing errors (in non-type safe languages like C)
 - aliasing violations ([TypeSanitizer](#))
 - mismatched types ([libcrunch](#))

Spec taxonomy (2)

- Parallel programming errors (Tsan):
 - deadlocks and data races
- Language-specific errors:
 - integer overflows ([UBsan](#))
 - static init order fiasco ([Asan](#))

Spec taxonomy (3)

- Invalid usage of APIs:
 - not checking return codes of syscalls or standard APIs
 - mismatched memory allocation API (calling `free` on `new`-ed pointer)
 - invalid comparators

Spec taxonomy (4)

- Violation of "business rules":
 - very application specific
 - specifications are extracted from domain experts, architects, QA, etc.
 - [Runtime Verification, from Theory to Practice and Back](#) and [Industrial Experiences with Runtime Verification](#)

Creating new checkers: instrumentation

The "instrumentation" part:

- for an existing spec, develop new ways to detect more errors more efficiently
- often determine whether checker will be used
- e.g. there were many buffer overflow checkers before [AddressSanitizer](#) but too slow or with limited coverage

Instrumentation taxonomy

- Aka [aspect-oriented programming](#) (AOP)
- Runtime verification is trivial in languages like Python or Java
 - full access to AST at runtime
 - many AOP frameworks
- Instrumentations of native code are categorized by stage in compilation pipeline and mechanism used for instrumentation
 - compromise between simplicity of implementation/integration and desired level of detail

Instrumentation taxonomy: preprocess-time

- Code can be instrumented by forced inclusion of debug header
 - e.g. via `-include mychecker.h`
 - header would contain something like

```
#define malloc my_safe_malloc
```
- Examples:
 - [dmalloc](#)
 - [Glibc FORTIFY SOURCE](#)

Instrumentation taxonomy: compile-time

- Compile-time instrumentation:
 - source-to-source (e.g. [libcrunch](#))
 - traditionally done via [CIL](#) but it's C only :(
 - [Clang LibTooling](#) supports C++ but is complicated to use due to baroque AST
 - codegen-based (e.g. [Asan](#) or [DirtyPad](#))
 - asm-based (e.g. [AFL](#) or [DirtyFrame](#))

Instrumentation taxonomy: link-time

Link-time instrumentation:

- replacing normal code with "checking" implementations at link time
- e.g. via `-Wl,--defsym,malloc=my_safe_malloc` or `-Wl,--wrap=malloc`
- e.g. [malloc dbg](#) replaces normal `malloc` if user links against debug version of Microsoft runtime

Instrumentation taxonomy: run-time

Run-time instrumentation types:

- `LD_PRELOAD`-based (e.g. [ElectricFence](#), [sortchecker](#), [failing-malloc](#))
 - `LD_PRELOAD` is a canonical way to implement AOP on Linux
- syscall instrumentation (e.g. [SystemTap](#))
- dynamic binary instrumentation (aka DBI, e.g. [Valgrind](#), [DynamoRIO](#) or [Intel Pin](#))

Creating new checkers: test corpus

- Project test suites provide insufficient code coverage
- Need to extend test corpus by generating new tests:
 - via fuzzing:
 - random (e.g. [Radamsa](#), [zzuf](#))
 - feedback-driven (e.g. [AFL](#), [libFuzzer](#))
 - concolic (e.g. [Microsoft SAGE](#), [Mayhem](#), [KLEE](#))
 - by developing generator for sufficiently important class of data
 - e.g. [Defensics](#) supports grammar-based test generation for [250+ protocols](#)
 - e.g. [Csmith](#) generates random C++ code for compiler testing

How to test a checker

- Once checker is ready you'll want to test it on as much code as you can
- Try to apply it to important OSS projects
 - archivers, media processing libraries, browsers, etc.
 - to find interesting package faster:
 - [package popularity rating](#)
 - [Debian codesearch](#) (supports both web and cmdline interfaces)

How to test an LD_PRELOAD- or DBI-based checker

- Checkers which do not require program recompilation are easier to test:
 - Run all apps in `/bin` and `/usr/bin`
 - without params, with `--help`, with `--version`
 - automatic but coverage is low (tests initialization code, at best)
 - boot complete Linux distro with your checker preloaded
 - for example [valgrind-preload](#)
 - limited applicability
 - need to perform manual actions to explore system behavior

How to test an arbitrary checker

- Run package unittests (if available)
 - good coverage but not scalable (5-30 minutes per package)
 - tiresome and demotivating :(
- System testsuites
 - run system benchmarks (e.g. [Phoronix suite](#) or [browser testsuites](#))
- Instrument complete Linux distro (e.g. [sanitize Tizen](#))
 - extremely hard...

How to test a checker: comparison

Test	Automatic	Coverage	All checkers
Running apps with standard params	Y	Low	Only LD_PRELOAD/DBI
System testsuites	Y	Average	Y
Manual package testing	N	High	Y
Distro boot	iff LD_PRELOAD/DBI	Average (need manual actions to increase)	Y

Using distro build systems

- Linux distros come with a vast number of packages
- Distro build systems can be reused
 - to apply checkers under the hood
 - and run package-specific unittests
- [Debian build toolchain](#)
 - Sadly only builds, not tests, but ...

debian_pkg_test

- With some hacking we can make Debian build system to run unittests for us!
- [debian_pkg_test](#) project
 - based on [pbuilder](#)
 - runs `make check` (or other standard test commands) once package build completes

Trends (1)

Increasing fuzzing speed and efficiency (coverage) by various means

- feedback-driven ("grey-box")
 - [AFL](#) and related tools (gofuzz, libfuzzer, etc.)
- symexec-driven ("white-box")
 - [Billions and Billions of Constraints: Whitebox Fuzz Testing in Production](#)
- various combinations thereof

Trends (2)

Increasing fuzzing adoption in community:

- integration of fuzzers into development lifecycles (kudos to @msh_smlv)
- inspire project owners to write fuzzing for their projects through initiatives like [OSS-fuzz](#)
- bug bounty programs e.g. [Google Fuzzilli](#)

Links

- [Runtime Verification conference](#) ([Springer](#))
 - Too scientific
 - Most papers are on verifying temporal logic assertions at runtime
- More practical: vulnerability reports
 - [CVE reports](#)
 - [DEFCON](#)
 - [Blackhat](#)
 - [Phrack](#)

Advertisement

Samsung System-On-Chip team is hiring developers to develop state-of-the-art compilers in Moscow Research Center:

- NPU Compiler Developer for Exynos AI Accelerator (<https://hh.ru/vacancy/42341825>)

Also need GPU performance engineers, whatever that means (<https://hh.ru/vacancy/44907512>).

The End

Please share your ideas on runtime verification (new checkers, novel ways to test them, etc.):

- tetra2005 beim gmail punct com
- TG [the_real_yugr](#)
- GH [.yugr](#)

Checker gotchas

- Instead of testing that bad objects are not accessed, make sure that such accesses cause havoc
 - Fill undef memory/regs with garbage (MSVS does this with malloced memory)
 - Unmap page after buffer to force segfault (ElectricFence)
 - Fill gaps in stack frame with random values (DirtyFrame)
 - Fill struct pads with random values (DirtyPad)
 - Intro random delays in Pthread-based programs